

# CS229: Machine Learning

Junchuan Zhao

May 24, 2022

## Abstract

This document includes my notes of CS229: Machine Learning.

## 1 Lecture 2

### 1.1 Linear Regression

Hypothesis:  $h(x) = \sum_{j=0}^n \theta_j x_j$  ( $n$ : number of features)

Linear Regression: the hypothesis is the linear combination of the training dataset.

Loss function (goal):  $\min_{\theta} \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)^2$

Parameters:  $\theta, m, n, x, y$

### 1.2 Gradient Descent

#### 1.2.1 Batch Gradient Descent

Basic Ideas: start with some  $\theta$ , keep changing  $\theta$  to reduce  $J(\theta)$ , repeat until convergent

$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$  ( $\alpha$ : learning rate)

$\frac{\partial}{\partial \theta_j} J(\theta) = (h_{\theta}(x) - y) \cdot x_j$

$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x)^i - y^i) \cdot x_j^i$  ( $\alpha$ : learning rate)

Batch Gradient Descent: all the training data as a batch when optimizing the loss function. (-): not good for large scale dataset, very expensive.

#### 1.2.2 Stochastic Gradient Descent

---

**Algorithm 1** Stochastic Gradient Descent

---

**Require:** the parameter of  $j^{th}$  feature

**Ensure:** the updated parameter of  $j^{th}$  feature

```
1: for  $i = 1$  to  $m$  do  
2:    $\theta_j := \theta_j - \alpha \cdot (h_{\theta}(x)^i - y^i) \cdot x_j^i$ ;  
3: end for
```

---

stochastic gradient descent heads over to the global optimum.

### 1.2.3 Total Equations

A Total Equation for Stochastic Gradient Descent:  $\nabla_{\theta} J(\theta) = \mathbf{0}$

If A is square ( $A \in \mathbb{R}_{n \times n}$ )

$$\text{Trace of A: } \text{tr}(A) = \sum_i A_{ii}$$

Features of Trace:

- $\text{tr}(A) = \text{tr}(A^T)$
- $f(A) = \text{tr}(AB), \nabla_A f(A) = B^T$
- $\text{tr}(AB) = \text{tr}(BA)$
- $\text{tr}(ABC) = \text{tr}(CAB)$
- $\nabla_A \text{tr}(AA^T C) = CA + C^T A$

$$\begin{aligned} \text{Loss Function: } \nabla_{\theta} J(\theta) &= \frac{1}{2}(X\theta - y)^T(X\theta - y) \\ &= \frac{1}{2}(\theta^T X^T - y^T)(X\theta - y) \\ &= \frac{1}{2}(\theta^T X^T X\theta - \theta^T X^T y - y^T X\theta) \\ &= X^T X\theta - X^T y \end{aligned}$$

$$\begin{aligned} \text{Loss Function: } X^T X\theta - X^T y &= \mathbf{0} \iff X^T X\theta = X^T y \\ \theta &= (X^T X)^{-1} X^T y \end{aligned}$$

## 2 Lecture 3

### 2.1 Locally Weighted Regression

"Parametric" learning algorithm: fit fixed set of parameters ( $\theta_i$ ) to data.

"Non-parametric learning algorithm": amount of data/parameters you need to keep grows (linearly) with the size of data.

Linear Regression: to evaluate  $h$  at certain  $x$

Fit  $\theta$  to minimize

$$\frac{1}{2} \sum_i (y^i - \theta^T x^i)^2,$$

return  $\theta^T x$

Linear Regression: to evaluate  $h$  at local region of  $x$

Fit  $\theta$  to minimize

$$\sum_i w^i (y^i - \theta^T x^i)^2, \text{ where } w^i \text{ is a "weight function".}$$

common choice for  $w^i$  is  $w^i = e^{-\frac{(x^i - x)^2}{2\tau^2}}$

If  $|x^i - x|$  is small, then  $w^i \approx 1$ .

If  $|x^i - x|$  is large, then  $w^i \approx 0$ .

$\tau$ : bandwidth, control a larger or narrower window.

### 2.2 Why Square Error?

Assume  $y^i = \theta^T x^i + \epsilon^i$ , where  $\epsilon^i \sim N(0, \sigma^2)$  models effects of random noise

$$\begin{aligned} p(y^i | x^i; \theta) &= \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(y^i - \theta^T x^i)^2}{2\sigma^2}} \\ \iff y^i | x^i; \theta &\sim N(\theta^T x^i, \sigma^2) \end{aligned}$$

Difference between Likelihood ( $L$ ) and Probability ( $P$ ):  $L$  varies parameters,  $P$  varies datapoints. Assume the datapoints are IID,

The "likelihood" of  $\theta$ :  $L(\theta) = p(\mathbf{y}|\mathbf{x}; \theta) = \prod_{i=1}^m p(y^i|x^i; \theta)$

$$l(\theta) = \log \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} e(\dots) = \sum_{i=1}^m [\log \frac{1}{\sqrt{2\pi}\sigma} + \log e(\dots)]$$

$$= m \log \frac{1}{\sqrt{2\pi}\sigma} - \sum_{i=1}^m \frac{(y^i - \theta^T x^i)^2}{2\sigma^2}$$

MLE: maximum likelihood estimation. Choose  $\theta$  to maximize  $L(\theta)$

i.e. choose  $\theta$  to minimize  $\frac{1}{2} \sum_{i=1}^m (y^i - \theta^T x^i)^2$ , which is actually  $J(\theta)$

## 2.3 Classification

Binary Classification:  $y \in \{0, 1\}$

### 2.3.1 Logistic Regression

We want  $h_\theta(x) \in [0, 1]$ , so we define

$$h_\theta(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

Sigmoid Function:  $g(z) = \frac{1}{1 + e^{-z}}$

An important thing is that: the partial derivative of the sigmoid function is a **concave** function, which means it has the global maximum.

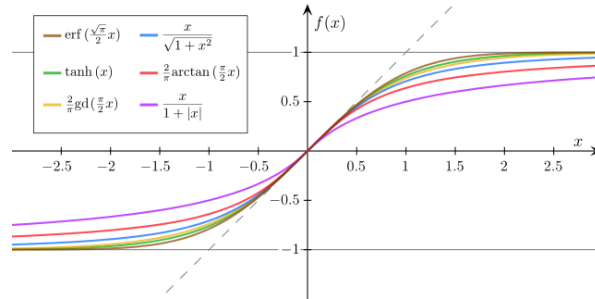


Figure 1: The graph of sigmoid function.

Different from Linear Regression  $h_\theta(x) = \theta^T x$ , Logistic Regression is  $h_\theta(x) = \frac{1}{1 + e^{-\theta^T x}}$ .

$$p(y=1|x; \theta) = h_\theta(x), p(y=0|x; \theta) = 1 - h_\theta(x)$$

$$\text{In sum, } p(y|x; \theta) = h_\theta(x)^y (1 - h_\theta(x))^{1-y}$$

$$L(\theta) = p(\mathbf{y}|\mathbf{x}; \theta) = \prod_{i=1}^m h_\theta(x^i)^{y^i} (1 - h_\theta(x^i))^{1-y^i}$$

$$l(\theta) = \log(L(\theta)) = \sum_{i=1}^m y^i \log h_\theta(x^i) + (1 - y^i) \log (1 - h_\theta(x^i))$$

Choose  $\theta$  to maximize  $l(\theta)$ , we use Batch gradient ascent.

Batch Gradient Ascent:

$$\theta_j := \theta_j + \alpha \frac{\partial}{\partial \theta_j} l(\theta)$$

$\theta_j := \theta_j + \alpha \sum_{i=1}^m (y^i - h_\theta(x^i)) x_j^i$ , same as the one in linear regression. (A larger category called generalized linear model (GLM))

Note: no **normal equation** for logistic regression solution.

### 2.3.2 Newton's Method

The basic idea: have some function  $f$ , we want to fit  $\theta$ , s.t.  $f(\theta) = 0 \iff$  want maximizing  $l(\theta)$   
 $\iff$  want  $l'(\theta) = 0$

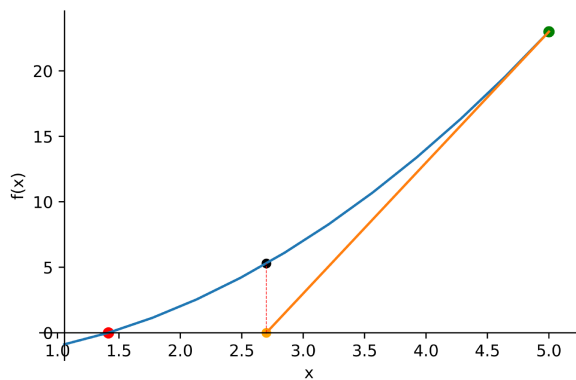


Figure 2: Newton's Method.

$$\theta^{t+1} := \theta^t - \frac{f(\theta^t)}{f'(\theta^t)}, \text{ let } f(\theta) = l'(\theta)$$

$$\theta^{t+1} := \theta^t - \frac{l'(\theta^t)}{l''(\theta^t)}$$

"Quadratic convergent": 0.01 error  $\longrightarrow$  0.0001 error  $\longrightarrow$  0.00000001 error.

When  $\theta$  is a vector: ( $\theta \in \mathbb{R}^{n+1}$ )

$$\theta^{t+1} := \theta^t + \alpha H^{-1} \nabla_{\theta} l(\theta)$$

where

$$\nabla_{\theta} l(\theta) = \begin{bmatrix} \frac{\partial l}{\partial \theta_{00}} & \frac{\partial l}{\partial \theta_{01}} & \cdots & \frac{\partial l}{\partial \theta_{0n}} \\ \frac{\partial l}{\partial \theta_{10}} & \frac{\partial l}{\partial \theta_{11}} & \cdots & \frac{\partial l}{\partial \theta_{1n}} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial l}{\partial \theta_{n0}} & \frac{\partial l}{\partial \theta_{n1}} & \cdots & \frac{\partial l}{\partial \theta_{nn}} \end{bmatrix}$$

$H \in \mathbb{R}^{n+1 \times n+1}$  is the Hessian Matrix,  $H_{ij} = \frac{\partial^2 l}{\partial \theta_i \partial \theta_j}$

Note: If the dataset is large, the size of  $H$  will be too large to compute.

## 3 Lecture 4

### 3.1 Perceptron

Binary step function:

$$g(z) = \begin{cases} 1 & z \geq 0 \\ 0 & z < 0 \end{cases} \quad (1)$$

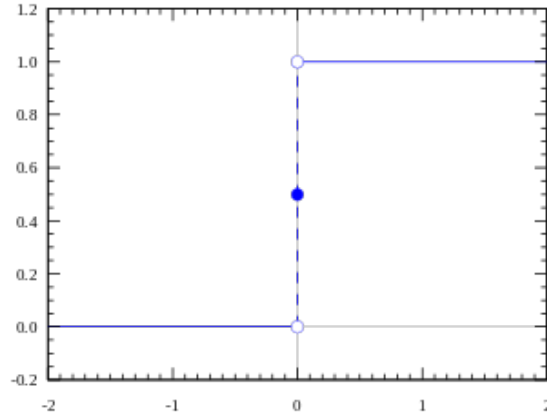


Figure 3: The graph of binary step.

Batch Gradient Ascent:  $\theta_j := \theta_j + \alpha(y^i - h_\theta(x^i))x_j^i$

## 3.2 Exponential Family

### 3.2.1 Defination and Examples of Exponential Family

Probability Density Function (PDF):  $p(y; \eta) = b(y)\exp[\eta^T T(y) - a(\eta)]$

$y$ : data

$\eta$ : natural parameter

$T(y)$ : sufficient statistic

$b(y)$ : base measure

$a(\eta)$ : log partition

Example 1: Bernoulli (Binary Data):  $p(y; \phi) = \phi^y(1 - \phi)^{(1-y)}$

$\phi$  = probability of the event

$$p(y; \phi) = \exp[y \log(\frac{\phi}{1-\phi}) + \log(1 - \phi)]$$

$$b(y) = 1$$

$$T(y) = y$$

$$\eta = \log \frac{\phi}{1-\phi} \Rightarrow \phi = \frac{1}{1+e^{-\eta}}$$

$$a(\eta) = -\log(1 - \phi) = -\log(1 - \frac{1}{1+e^{-\eta}}) = \log(1 + e^\eta)$$

Example 2: Gaussian (assume variance = 1):  $p(y; \mu) = \frac{1}{\sqrt{2\pi}}e^{-(\frac{y-\mu}{2})^2} = \frac{1}{\sqrt{2\pi}}e^{-\frac{y^2}{2}}\exp(\mu y - \frac{\mu^2}{2})$

$$b(y) = \frac{1}{\sqrt{2\pi}}e^{-\frac{y^2}{2}}$$

$$T(y) = y$$

$$\eta = \mu$$

$$a(\eta) = \frac{\mu^2}{2} = \frac{\eta^2}{2}$$

### 3.2.2 Properties of Exponential Family

- MLE w.r.t  $\eta$  is concave, NLL is convex
- $E[y; \eta] = \frac{\partial}{\partial \eta} a(\eta)$
- $Var[y; \eta] = \frac{\partial^2}{\partial \eta^2} a(\eta)$

Real - Gaussian; Binary - Bernoulli; Count - Poisson;  $\mathbb{R}^+$  - Gamma, Exponential; Distribution - Beta, Dirichlet (Bayesian).

### 3.3 Generalized Linear Model (GLM)

#### 3.3.1 GLM Assumptions

Assumptions/Design Choices

- (1)  $y|x; \theta \sim \text{Exponential Family } (\eta)$
- (2)  $\eta = \theta^T x$ ,  $\theta \in \mathbb{R}^n$  and  $x \in \mathbb{R}^n$
- (3) Test Time: output  $h_\theta(x) = E[y|x; \theta]$

#### 3.3.2 GLM Pipeline

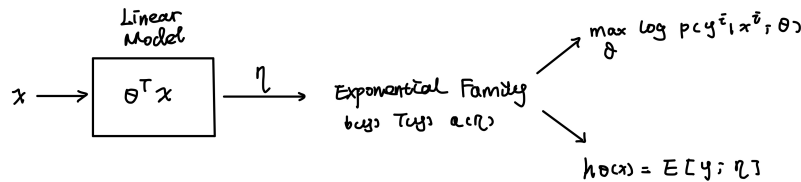


Figure 4: The pipeline for GLM.

#### 3.3.3 GLM Training

Learning Update Rule:  $\theta_j := \theta_j + \alpha(y^i - h_\theta(x^i))x_j^i$

Terminology:

$\eta$  - natural parameter

$g(\eta) = \mu = E(y; \eta)$  - canonical response function

$\eta = g^{-1}(\mu)$  - canonical link function

Three parameterizations:

Model Param:  $\theta$  (**Learn Parameters**)

Natural Param:  $\eta$  (Design Choice:  $\theta^T x$ )

Canonical Param:  $\phi$  - Bernoulli,  $\mu \sigma$  - Gaussian,  $\lambda$  - Poisson (Canonical Response:  $g(\cdot)$ )

#### 3.3.4 Review Logistic Regression

$h_\theta(x) = E[y|x; \theta] = \phi$  (the mean of the Bernoulli distribution)

$$h_\theta(x) = \phi = \frac{1}{1+e^{-\eta}} = \frac{1}{1+e^{-\theta^T x}}$$

### 3.4 Softmax Regression

Task: multi-class Classification

$k$  - number of classes,  $x^i \in \mathbb{R}^n$ , Label  $y^i \in \{0, 1\}^k$  (one-hot vector)

Each class has its own parameters:  $\theta_{class} \in \mathbb{R}^n$

$$p(y = i|x; \theta) = \frac{e^{\theta_i^T x}}{\sum_{j=1}^k e^{\theta_j^T x}}$$

$$l(\theta) = \sum_{i=1}^m \log \prod_{l=1}^k \left( \frac{e^{\theta_l^T x^i}}{\sum_{j=1}^k e^{\theta_j^T x^i}} \right)^{\{y^i=l\}}$$

## 4 Lecture 5

### 4.1 Discriminative v.s. Generative Learning Algorithms

Discriminative: Learn  $p(y|x)$  which is the mapping from  $x$  to  $y$ .

Generative: Learn  $p(x|y)$ ,  $p(y)$  (class prior)

According to Bayes rule:  $p(y|x) = \frac{p(x|y)p(y)}{p(x)}$ ,  $p(x) = \sum_{i=0}^k p(x|y=i)p(y=i)$

### 4.2 Gaussian Discriminant Analysis (GDA)

#### 4.2.1 Assumptions and Fundamental Knowledge

Suppose  $x \in \mathbb{R}^n$  (convention: drop  $x_0 = 1$ )

Assume  $p(x|y)$  is Gaussian

Basic Knowledge for Multivariate Gaussian:

$$z \sim N(\mu, \Sigma), (z \in \mathbb{R}^n)$$

$$p(z) = \frac{1}{(2\pi)^{\frac{n}{2}} |\Sigma|^{\frac{1}{2}}} \exp(-\frac{1}{2}(z - \mu)^T \Sigma^{-1} (z - \mu)) \quad E[z] = \mu$$

$$Cov(z) = E[(z - \mu)(z - \mu)^T] = E_{zz^T} - E[z]E[z]^T$$

#### 4.2.2 GDA model

Parameters:  $\mu_0, \mu_1, \Sigma$  (use same covariance matrix)

$$p(x|y=0) = \frac{1}{(2\pi)^{\frac{n}{2}} |\Sigma|^{\frac{1}{2}}} \exp(-\frac{1}{2}(x - \mu_0)^T \Sigma^{-1} (x - \mu_0))$$

$$p(x|y=1) = \frac{1}{(2\pi)^{\frac{n}{2}} |\Sigma|^{\frac{1}{2}}} \exp(-\frac{1}{2}(x - \mu_1)^T \Sigma^{-1} (x - \mu_1))$$

$$p(y) = \phi^y (1 - \phi)^{1-y} \quad (p(y=1) = \phi)$$

Training set:  $\{(x^i, y^i)\}_{i=1}^m$

$$\text{Joint Likelihood: } L(\phi, \mu_0, \mu_1, \Sigma) = \prod_{i=1}^m p(x^i, y^i; \phi, \mu_0, \mu_1) = \prod_{i=1}^m p(x^i|y^i)p(y^i)$$

$$\text{Discriminative Learning (Conditional Likelihood): } \prod_{i=1}^m p(y^i|x^i; \theta)$$

Maximize Likelihood Estimation (MLE):

$$\max_{\phi, \mu_0, \mu_1, \Sigma} l(\phi, \mu_0, \mu_1, \Sigma)$$

Results of MLE:

$$\phi = \frac{\sum_{i=1}^m y^i}{m} = \frac{\sum_{i=1}^m \mathbb{1}\{y^i=1\}}{m} \quad (\mathbb{1} : \mathbb{1}\{\text{true}\} = 1, \mathbb{1}\{\text{false}\} = 0)$$

$$\mu_0 = \frac{\sum_{i=1}^m \mathbb{1}\{y^i=0\} x^i}{\sum_{i=1}^m \mathbb{1}\{y^i=0\}} = \frac{\text{sum of feature vectors for samples with } y=0}{\text{number of samples with } y=0}$$

$$\mu_1 = \frac{\sum_{i=1}^m \mathbb{1}\{y^i=1\} x^i}{\sum_{i=1}^m \mathbb{1}\{y^i=1\}}$$

$$\Sigma = \frac{1}{m} \sum_{i=1}^m (x^i - \mu_{y^i})(x^i - \mu_{y^i})^T$$

Prediction:

$$\begin{aligned} \arg \max_y p(y|x) &= \arg \max_y \frac{p(x|y)p(y)}{p(x)} \quad (p(x) \text{ is constant}) \\ &= \arg \max_y p(x|y)p(y) \end{aligned}$$

### 4.2.3 Compare GDA to Logistic Regression

For fix parameters:  $\phi, \mu_0, \mu_1, \Sigma$ , plot  $p(y|x; \phi, \mu_0, \mu_1, \Sigma)$  as a function of  $x$ .

By Bayes Rule:  $p(y = 1|x; \phi, \mu_0, \mu_1, \Sigma) = \frac{p(x|y=1; \mu_0, \mu_1, \Sigma)p(y=1; \phi)}{p(x; \phi, \mu_0, \mu_1, \Sigma)}$ .

In the end,  $p(y = 1|x; \phi, \mu_0, \mu_1, \Sigma) = \frac{1}{1+e^{xp(-\theta^T x)}}$ , where  $\theta$  is a function of  $\phi, \mu_0, \mu_1, \Sigma$ .

The assumption of GDA  $\Rightarrow p(y = i|x) = \frac{1}{1+e^{-\theta^T x}}$

However,  $p(y = i|x) = \frac{1}{1+e^{-\theta^T x}} \not\Rightarrow$  GDA assumptions.

| GDA   | Logistic Regression                         |
|---|---|
| stronger assumptions                        | weaker assumptions                          |
| perform better when assumptions are correct | fine with all the scenarios (distributions) |
| small dataset                               | large dataset                               |
| computationally efficient                   | fits the era of the big data                |

Table 1: The comparison between GDA and Logistic Regression.

### 4.3 Naive Bayes

For a spam email classification problem:

1. create a word dictionary (size  $m$ ).

If represent  $x$  by a binary feature vector,  $x \in \{0, 1\}^n$ . ( $x_i = \mathbb{1}\{\text{word } i \text{ appears in email}\}$ )

This cause  $2^n$  possible values of  $x$ , which is not possible for large dictionary.

2. Assume  $x_i$ 's are **conditionally independent given  $y$** , which means

$$p(x_1, \dots, x_{10000}|y) = p(x_1|y)p(x_2|x_1, y)p(x_3|x_1, x_2, y) \dots p(x_{10000}|x_1, \dots, x_{9999}, y)$$

$$\stackrel{\text{assume}}{=} p(x_1|y)p(x_2|y) \dots p(x_{10000}|y) = \prod_{i=1}^n p(x_i|y)$$

Parameters:

$$\phi_{j|y=1} = p(x_j = 1|y = 1)$$

$$\phi_{j|y=0} = p(x_j = 1|y = 0)$$

$$\phi_y = p(y = 1)$$

$$\text{Joint Likelihood: } L(\phi_i, \phi_{j|y}) = \prod_{i=1}^m p(x^i, y^i; \phi_j, \phi_{j|y})$$

$$\text{MLE results: } \phi_y = \frac{\sum_{i=1}^m \mathbb{1}\{y^i=1\}}{m}$$

$$\phi_{j|y=1} = \frac{\sum_{i=1}^m \mathbb{1}\{x_j^i=1, y^i=1\}}{\sum_{i=1}^m \mathbb{1}\{y^i=1\}}$$

$$\phi_{j|y=0} = \frac{\sum_{i=1}^m \mathbb{1}\{x_j^i=1, y^i=0\}}{\sum_{i=1}^m \mathbb{1}\{y^i=0\}}$$

At prediction time:

$$p(y = 1|x) = \frac{p(x|y=1)p(y=1)}{p(x|y=1)p(y=1) + p(x|y=0)p(y=0)}$$

Problem: when a word does not exist -  $\frac{0}{0}$



## 5 Lecture 6

### 5.1 Laplace Smoothing

$$x \in \{1, \dots, k\}, \text{ estimate } p(x = j) = \frac{\sum_{i=1}^m \mathbb{1}\{x^i = j\} + 1}{m + k}$$

$$\phi_{j|y=0} = \frac{\sum_{i=1}^m \mathbb{1}\{x_j^i = 1, y^i = 0\} + 1}{\sum_{i=1}^m \mathbb{1}\{y^i = 0\} + 2}$$

### 5.2 Multinomial Event Model

$$x \in \begin{bmatrix} 600 \\ 800 \\ 1600 \\ 6200 \end{bmatrix}, x \in \mathbb{R}^n$$

$x_j \in \{1, 2, \dots, 10000\}$ ,  $n$  = length of email,  $m$  = number of emails

Parameters:  $\phi_y = p(y = 1)$ ,  $\phi_{k|y=0} = p(x_j = k|y = 0)$ ,  $\phi_{k|y=1} = p(x_j = k|y = 1)$

$$\text{MLE results: } \phi_{k|y=0} = \frac{\sum_{i=1}^m (\mathbb{1}\{y^i = 0\} \sum_{j=1}^{n_i} \mathbb{1}\{x_j^i = k\})}{\sum_{i=1}^m \mathbb{1}\{y^i = 0\} n_i}$$

$$\underline{\underline{\text{laplace smoothing}}} = \frac{\sum_{i=1}^m (\mathbb{1}\{y^i = 0\} \sum_{j=1}^{n_i} \mathbb{1}\{x_j^i = k\}) + 1}{\sum_{i=1}^m \mathbb{1}\{y^i = 0\} n_i + 10000}$$

The advantages for GDA and Naive Bayes: quick to train, non-iterative.

### 5.3 Support Vector Machines SVM

SVM helps to find non-linear boundaries for classification problems.

#### 5.3.1 Optimal Margin Classifier

Functional Margin: how confidently and accurately to classify an example.

Logistic Regression:  $h_\theta(x) = g(\theta^T x)$ , predict 1 if  $\theta^T x \geq 0$  and 0 otherwise.

We want if  $y^i = 1$ , hope that  $\theta^T x^i \gg 0$  and 0 when  $\theta^T x^i \ll 0$ . – correct or confident prediction.

Notation: Labels  $y \in \{-1, +1\}$ ,

$$g(z) = \begin{cases} 1 & z \geq 0 \\ -1 & z < 0 \end{cases} \quad (2)$$

$h_{W,b}(x) = g(W^T x + b)$ ,  $W \in \mathbb{R}^n$  and  $b \in \mathbb{R}$ . (Compare to Logistic Regression,  $b$  is  $\theta_0$ )

Functional Margin of hyperplane defined by  $(W, b)$  w.r.t  $(x^i, y^i)$ :

$\hat{\gamma}^i = y^i(W^T x^i + b)$ , we want  $\hat{\gamma}^i \gg 0$

Functional Margin w.r.t training set:  $\hat{\gamma} = \min_i \hat{\gamma}^i$ ,  $i = 1, \dots, m$

Rescaling the parameters:  $(W, b) \rightarrow (\frac{W}{\|W\|}, \frac{b}{\|W\|})$

Geometric Margin: the distance from the data point to the hyperplane.

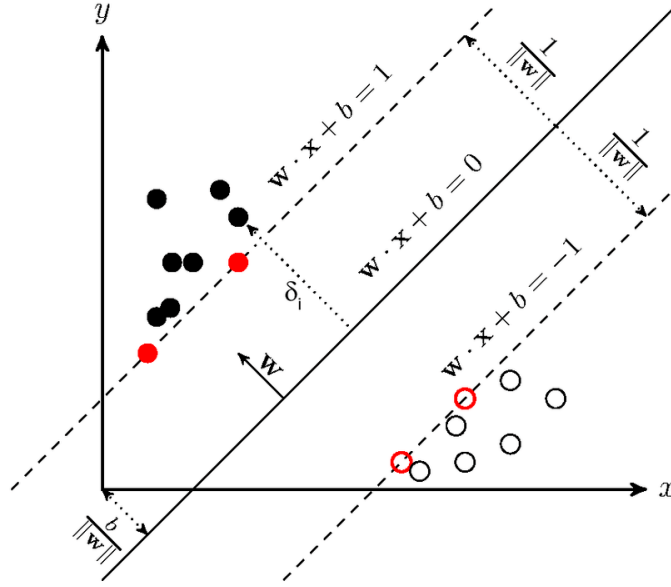


Figure 5: Geometric Margin.

Geometric Margin of hyperplane  $(W, b)$  w.r.t  $(x^i, y^i)$ :

$$\gamma^i = \frac{y^i(W^T x^i + b)}{\|W\|}$$

Geometric Margin w.r.t training set:  $\gamma = \min_i \gamma^i$

The relationship between geometric and functional margin:  $\gamma^i = \frac{\hat{\gamma}^i}{\|W\|}$

$\hat{\gamma}$  is functional margin;  $\gamma$  is geometric margin.

Optimal margin classifier: choose  $w, b$  to maximize  $\gamma$ :

$$\max_{\gamma, w, b} \gamma \text{ s.t. } \frac{y^i(W^T x^i + b)}{\|W\|} \geq \gamma, (i = 1, \dots, m.) \Rightarrow \min_{w, b} \|W\|^2 \text{ s.t. } y^i(W^T x^i + b) \geq 1$$

## 6 Lecture 7

### 6.1 Optimal Margin Classifiers

If  $x^i \in \mathbb{R}^{100}$  (highly dimensional spaces), suppose  $W = \sum_{i=1}^m \alpha_i y^i x^i$  ( $W$  can be represented by the linear combination of the training set)

Intuition 1: Logistic Regression  $\theta := 0$ , Gradient Descent:  $\theta := \theta - \alpha(h_\theta(x^i) - y^i)x^i$

$$\text{Batch Gradient Descent: } \theta := \theta - \alpha \sum_{i=1}^m (h_\theta(x^i) - y^i)x^i$$

Intuition 2:  $W$  pins the decision boundary,  $W$  lies in the span of the training examples

$$\begin{aligned} \min_{W, b} \frac{1}{2} \|W\|^2 \text{ s.t. } y^i(W^T x^i + b) &\geq 1, W = \sum_{i=1}^m \alpha_i y^i x^i \\ \min_{W, b} \frac{1}{2} \left( \sum_{i=1}^m \alpha_i y^i x^i \right)^T \left( \sum_{j=1}^m \alpha_j y^j x^j \right) &= \min_{W, b} \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y^i y^j \langle x^i, x^j \rangle \\ y^i(W^T x^i + b) &\geq 1 \Leftrightarrow y^i \left( \sum_j \alpha_j y^j \langle x^j, x^i \rangle + b \right) \geq 1 \end{aligned}$$

### 6.1.1 Dual Optimization Problem

$$\begin{aligned} \max_{\alpha} W(\alpha) &= \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^i y^j \alpha_i \alpha_j \langle x^i, x^j \rangle \\ \text{s.t. } \alpha_i &\geq 0 \ (i = 1, \dots, m), \sum_{i=1}^m \alpha_i y^i = 0 \end{aligned}$$

Make predictions:

(1) solve for  $\alpha_i$ 's

(2) to make prediction,  $h_{w,b}(x) = g(w^T x + b) = g((\sum_i \alpha_i y^i x^i)^T x + b) = g(\sum_i \alpha_i y^i \langle x^i, x \rangle + b)$

## 6.2 Kernel

### 6.2.1 Kernel Tricks

Kernel trick:

- (1) write algorithm in terms of  $\langle x, z \rangle$
- (2) map  $x \rightarrow \phi(x)$  (higher dimension)
- (3) find a way to compute  $K(x, z) = \phi(x)^T \phi(z)$
- (4) replace  $\langle x, z \rangle$  in algorithm with  $K(x, z)$

Traditional way:

$$x = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad \phi(x) = \begin{bmatrix} x_1 x_1 \\ x_1 x_2 \\ x_1 x_3 \\ \dots \\ x_3 x_2 \\ x_3 x_3 \end{bmatrix} \quad \phi(z) = \begin{bmatrix} z_1 z_1 \\ z_1 z_2 \\ z_1 z_3 \\ \dots \\ z_3 z_2 \\ z_3 z_3 \end{bmatrix}$$

$n^2$  elements need  $O(n^2)$  time to compute  $\phi(x)$  or  $\phi(x)^T \phi(z)$

Kernel trick # 1: ( $O(n)$ )

$$\begin{aligned} K(x, z) &= \phi(x)^T \phi(z) = (x^T z)^2 = \left( \sum_{i=1}^n x_i z_i \right) \left( \sum_{j=1}^n x_j z_j \right) \\ &= \sum_{i=1}^n \sum_{j=1}^n (x_i x_j) (z_i z_j) \end{aligned}$$

Kernel trick # 2:  $K(x, z) = (x^T z + c)^2$

Kernel trick # 3:  $K(x, z) = (x^T z + c)^d$ ,  $\phi(x)$  has all  $\binom{n+d}{d} \approx (n+d)^d$  features of monomials up to order  $d$ .

SVM = Optimal Margin Classifier + Kernel Tricks (Map the data points to a higher dimension, then separate by using linear hyperplane)

### 6.2.2 How to Make Kernels?

Basic idea: if  $x, z$  are similar,  $K(x, z) = \phi(x)^T \phi(z)$  is large. Vice versa.

Gaussian Kernel:  $K(x, z) = \exp(-\frac{\|x-z\|^2}{2\sigma^2})$

Theorem(Mercer):  $K$  is a valid kernel function (i.e.  $\exists \phi$  s.t.  $K(x, z) = \phi(x)^T \phi(z)$ )  $\iff$  for any  $d$  points  $(x^1 \dots x^d)$ , the corresponding kernel matrix  $K \geq 0$ .

Linear Kernel:  $K(x, z) = x^T z$ ,  $\phi(x) = x$

Gaussian Kernel:  $K(x, z) = \exp(-\frac{\|x-z\|^2}{2\sigma^2})$ ,  $\phi(x) \in \mathbb{R}^\infty$

Note: a lot of algorithms can marry with the kernel tricks to reduce the amounts of computation.

### 6.3 L1-norm Soft Margin SVM

$$\min_{W, b, \xi_i} \frac{1}{2} \|W\|^2 + C \sum_{i=1}^m \xi_i \text{ s.t. } y^i (W^T x^i + b) \geq 1 - \xi_i \quad (i = 1, \dots, m; \xi_i \geq 0)$$

Note: more robust to outliers.

Dual Optimization Form:

$$\max_{\alpha} W(\alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^i y^j \alpha_i \alpha_j \langle x^i, x^j \rangle$$

$$\text{s.t. } 0 \leq \alpha_i \leq C \quad (i = 1, \dots, m), \quad \sum_{i=1}^m \alpha_i y^i = 0$$

### 6.4 SVM Kernels

$$K(x, z) = (x^T z)^d = \exp(-\frac{\|x-z\|^2}{2\delta^2})$$

#### 6.4.1 Protean Sequence Classifier

Protean Sequence: *BAJTSJAIBAJT...*

Protean Codebook: AAAA, AAAB, AAAC, ..., ZZZZ ( $20^4 = 160000$ )

$$\phi(x) = \begin{bmatrix} 0 \\ 0 \\ 2 \\ 0 \\ \dots \\ 4 \\ \dots \\ 0 \end{bmatrix}$$

$$\phi(x)^T \phi(z) = K(x, z)$$

## 7 Lecture 8

### 7.1 Bias and Variance

Underfitting problem: high bias.

Overfitting problem: high variance.

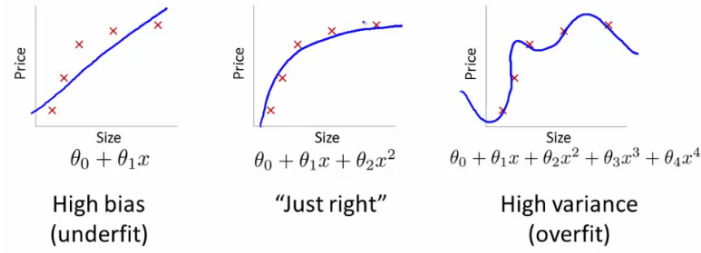


Figure 6: Bias and Variance.

## 7.2 Regularization

Regularization for linear regression:  $\min_{\theta} \frac{1}{2} \sum_{i=1}^m \|y^i - \theta^T x^i\|^2 + \frac{\lambda}{2} \|\theta\|^2$

When  $\lambda$  is too big – underfitting, whereas  $\lambda$  is too small – overfitting

Regularization for logistic regression:  $\arg \max_{\theta} \sum_{i=1}^m \log p(y^i | x^i; \theta) - \lambda \|\theta\|^2$

Support Vector Machine (SVM):  $\min_{W, b} \|W\|^2$ , since  $W$  is  $\theta_1 \dots \theta_n$ , the loss function already involves the regularization idea.

From another perspective:  $S : \{(x^i, y^i)\}_{i=1}^m$

$$p(\theta|S) = \frac{p(S|\theta)p(\theta)}{p(S)}$$

$$\arg \max_{\theta} p(\theta|S) = \arg \max_{\theta} p(S|\theta)p(\theta)$$

$$= \arg \max_{\theta} \left( \prod_{i=1}^m p(y^i | x^i, \theta) \right) p(\theta)$$

$$\text{Assume } \theta \sim N(0, \tau^2 I), p(\theta) = \frac{1}{\tau \sqrt{2\pi I}} \exp\left(-\frac{\theta^2}{2\tau^2 I}\right)$$

After log the MLE, we will obtain  $\|\theta\|^2$  term which is the regularization term.

Frequentist:  $\arg \max_{\theta} p(S|\theta)$  – MLE (estimate the parameters)

Bayesian:  $\theta$  is unknown, prior distribution  $p(\theta)$ ,  $\arg \max_{\theta} p(\theta|S)$  – MAP (maximum a posterior)

## 7.3 Data Split

Train/dev/test sets

Choose the form of polynomial,  $\lambda$ ,  $\tau$ ,  $C$

### 7.3.1 Hold-out Cross Validation (large dataset)

$S \rightarrow S_{train}, S_{dev}, S_{test}$  (dev: development)

- (1) Train each model (options for design of polynomial) on  $S_{train}$ , and get some hypothesis  $h$ .
- (2) Measure error on  $S_{dev}$ , pick model w.r.t lowest error on  $S_{dev}$ .
- (3) Measure error on  $S_{test}$  and report.

### 7.3.2 k-fold Cross Validation (small dataset)

$m = 100$ ,  $S = \{(x^i, y^i)\} (i = 1, \dots, m)$

$k = 10$  is typical,  $k = 5$  for illustration (Divide the dataset into 5 pieces).

---

**Algorithm 2** k-fold CV

---

```
1: for  $d = 1$  to 5 do
2:   for  $i = 1$  to  $k$  do
3:     train (fit parameters) on  $k - 1$  pieces and test on the training one piece;
4:   end for
5:   average;
6: end for
```

---

$k$ -fold Cross Validation makes more efficient use of data, whereas it is computationally very expensive.

## 7.4 Feature Selection

---

**Algorithm 3** Feature Selection

---

```
1: Start with  $\mathbb{F} = \phi$ ;
2: for  $i$  in feature bags do
3:   Try adding each feature  $i$  to  $\mathbb{F}$ , and see which single feature addition most improves the  $S_{dev}$  performance;
4:   Add the feature to  $\mathbb{F}$ ;
5: end for
```

---

Illustration:

$x_1, \dots, x_5$

(1)  $\phi$ : empty set of features,  $h(x) = \theta_0$

(2)  $\begin{bmatrix} \phi + x_1 \\ \phi + x_2 \\ \dots \\ \phi + x_5 \end{bmatrix} \begin{bmatrix} h_\theta(x) = \theta_0 + \theta_1 x_1 \\ h_\theta(x) = \theta_0 + \theta_1 x_2 \\ \dots \\ h_\theta(x) = \theta_0 + \theta_1 x_5 \end{bmatrix}$ ,  $x_2$  achieves the best performance,  $\mathbb{F} = \{x_2\}$ .

(3)  $\begin{bmatrix} x_2 + x_1 \\ x_2 + x_3 \\ x_2 + x_4 \\ x_2 + x_5 \end{bmatrix}$ ,  $\mathbb{F} = \{x_2, x_4\}$

...