

CS229: Machine Learning

Junchuan Zhao

May 17, 2022

Abstract

This document includes my notes of CS229: Machine Learning.

1 Lecture 2

1.1 Linear Regression

Hypothesis: $h(x) = \sum_{j=0}^n \theta_j x_j$ (n : number of features)

Linear Regression: the hypothesis is the linear combination of the training dataset.

Loss function (goal): $\min_{\theta} \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)^2$

Parameters: θ, m, n, x, y

1.2 Gradient Descent

1.2.1 Batch Gradient Descent

Basic Ideas: start with some θ , keep changing θ to reduce $J(\theta)$, repeat until convergent

$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$ (α : learning rate)

$\frac{\partial}{\partial \theta_j} J(\theta) = (h_{\theta}(x) - y) \cdot x_j$

$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x)^i - y^i) \cdot x_j^i$ (α : learning rate)

Batch Gradient Descent: all the training data as a batch when optimizing the loss function. (-): not good for large scale dataset, very expensive.

1.2.2 Stochastic Gradient Descent

Algorithm 1 Stochastic Gradient Descent

Require: the parameter of j^{th} feature

Ensure: the updated parameter of j^{th} feature

```
1: for  $i = 1$  to  $m$  do  
2:    $\theta_j := \theta_j - \alpha \cdot (h_{\theta}(x)^i - y^i) \cdot x_j^i$ ;  
3: end for
```

stochastic gradient descent heads over to the global optimum.

1.2.3 Total Equations

A Total Equation for Stochastic Gradient Descent: $\nabla_{\theta} J(\theta) = \mathbf{0}$

If A is square ($A \in \mathbb{R}_{n \times n}$)

$$\text{Trace of A: } tr(A) = \sum_i A_{ii}$$

Features of Trace:

- $tr(A) = tr(A^T)$
- $f(A) = tr(AB), \nabla_A f(A) = B^T$
- $tr(AB) = tr(BA)$
- $tr(ABC) = tr(CAB)$
- $\nabla_A tr(AA^T C) = CA + C^T A$

$$\begin{aligned} \text{Loss Function: } \nabla_{\theta} J(\theta) &= \frac{1}{2}(X\theta - y)^T(X\theta - y) \\ &= \frac{1}{2}(\theta^T X^T - y^T)(X\theta - y) \\ &= \frac{1}{2}(\theta^T X^T X\theta - \theta^T X^T y - y^T X\theta) \\ &= X^T X\theta - X^T y \end{aligned}$$

$$\begin{aligned} \text{Loss Function: } X^T X\theta - X^T y &= \mathbf{0} \iff X^T X\theta = X^T y \\ \theta &= (X^T X)^{-1} X^T y \end{aligned}$$

2 Lecture 3

2.1 Locally Weighted Regression

"Parametric" learning algorithm: fit fixed set of parameters (θ_i) to data.

"Non-parametric learning algorithm": amount of data/parameters you need to keep grows (linearly) with the size of data.

Linear Regression: to evaluate h at certain x

Fit θ to minimize

$$\frac{1}{2} \sum_i (y^i - \theta^T x^i)^2,$$

return $\theta^T x$

Linear Regression: to evaluate h at local region of x

Fit θ to minimize

$$\sum_i w^i (y^i - \theta^T x^i)^2, \text{ where } w^i \text{ is a "weight function".}$$

common choice for w^i is $w^i = e^{(-\frac{(x^i - x)^2}{2\tau^2})}$

If $|x^i - x|$ is small, then $w^i \approx 1$.

If $|x^i - x|$ is large, then $w^i \approx 0$.

τ : bandwidth, control a larger or narrower window.

2.2 Why Square Error?

Assume $y^i = \theta^T x^i + \epsilon^i$, where $\epsilon^i \sim N(0, \sigma^2)$ models effects of random noise

$$\begin{aligned} p(y^i | x^i; \theta) &= \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(y^i - \theta^T x^i)^2}{2\sigma^2}} \\ \iff y^i | x^i; \theta &\sim N(\theta^T x^i, \sigma^2) \end{aligned}$$

Difference between Likelihood (L) and Probability (P): L varies parameters, P varies datapoints.
 Assume the datapoints are IID,

The "likelihood" of θ : $L(\theta) = p(\mathbf{y}|\mathbf{x};\theta) = \prod_{i=1}^m p(y^i|x^i;\theta)$

$$\begin{aligned} l(\theta) &= \log \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} e^{(\dots)} = \sum_{i=1}^m [\log \frac{1}{\sqrt{2\pi}\sigma} + \log e^{(\dots)}] \\ &= m \log \frac{1}{\sqrt{2\pi}\sigma} - \sum_{i=1}^m \frac{(y^i - \theta^T x^i)^2}{2\sigma^2} \end{aligned}$$

MLE: maximum likelihood estimation. Choose θ to maximize $L(\theta)$

i.e. choose θ to minimize $\frac{1}{2} \sum_{i=1}^m (y^i - \theta^T x^i)^2$, which is actually $J(\theta)$